



TOKENS: Currency of LLM's



Tokens are stored into arrays of numbers.

⇒ LLM's split the input text into the longest known word/subword in their vocabulary.



⇒ LESS TOKEN = ↑ EFF.

⇒ MORE TOKENS = ↑ PROCESSING BY LLM'S

⇒ ↑ VOCAB. SIZE = ↓ PROCESSING POWER

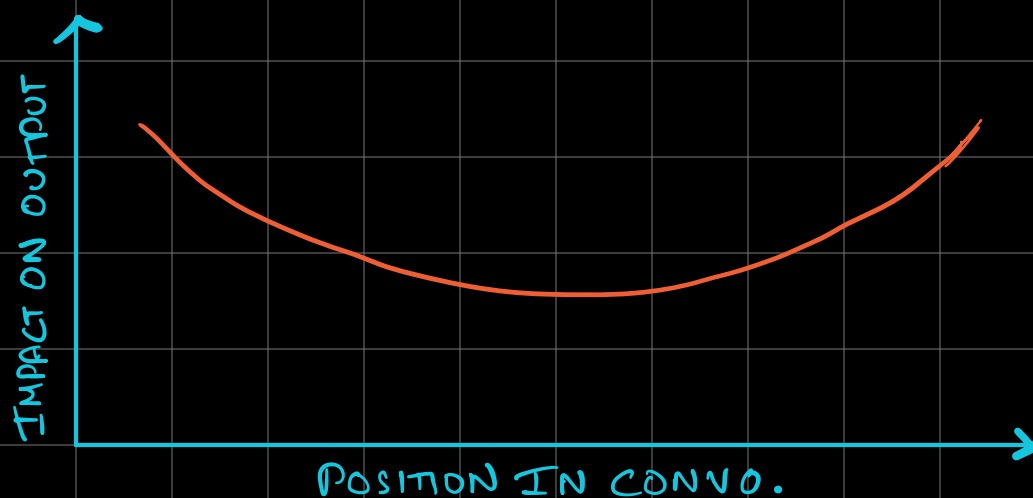
⇒ ↑ VOCAB. SIZE = ↑ MODEL

CONTEXT WINDOW: It is the hard arch. limit to the maximum no of tokens the model can process simultaneously.

CONTEXT: The sequence of tokens that a model is actively utilizing to predict the output.

CONTEXT WINDOW LIMITS: Each model has a limit that is hard coded by the dev of the model.

⇒ Larger the context window, more the model's performance degrades.



⇒ This is an emergent property due to the arch. of the model.